RAMCHANDRA REDDY CHINTALA

SENIOR DEVOPS ENGINEER

CONTACT

Hyderabad

076758 86590

ramcreddy.ch@gmail.com

PROFESSIONAL SUMMARY

With over 12 years of IT experience, I specialize in managing complex environments, implementing best DevOps and Cloud practices, and optimizing systems for peak efficiency. My expertise spans DevOps, Cloud, Hadoop, Kafka, Linux, and Middleware Management. Certified in AWS Solutions Architect, Microsoft Azure DevOps, Kubernetes Administration, Amazon EKS, and Cloudera Hadoop Administration, I am consistently recognized for delivering impactful solutions and staying current with industry developments. Excited to bring my skills and passion to innovative projects.

SKILLS

- Container Orchestration: OpenShift, Kubernetes
- Infrastructure as Code (IAAC) & Configuration Management (CM): Terraform, Ansible, GitOps
- Continuous Integration (CI) Tools: Jenkins, Azure DevOps, Bamboo, AWS CodeBuild
- DevSecOps: OWASP Dependency-Check, Clair, Terrascan, SonarQube, HashiCorp Vault
- Big Data, Messaging & Streaming: Kafka, Confluent Platform, Hadoop, NiFi
- ELK Stack: Elasticsearch, Kibana, Filebeat, Logstash
- Programming Languages & Scripting: Shell Scripting, Python
- Project Management & Collaboration: Jira, Remedy, ServiceNow, Atlassian Stack, Artifactory
- Logging & Analysis: Datadog, Datadog, Grafana, Splunk,

EXPERIENCE

Senior DevOps Engineer SagarSoft, Hyderabad | Client:Point72, USA
10/2017 - Present

- Administered Kubernetes and OpenShift clusters, implementing RBAC controls for a 20% vulnerability reduction, auto-scaling for 95% uptime, and optimizing resource utilization by 15%.
- Spearheaded the automated migration of 300+ apps from OpenShift 3.11 to 4.9, ensuring zero downtime, minimizing errors, and saving over 168 hours of development time.
- Architected CI/CD pipelines with Jenkins and OpenShift, automating builds for 500+ microservices, reducing release cycles from 2 weeks to 24 hours.
- Architected and deployed high-availability AWS applications with 99.9% uptime, reducing infrastructure costs by 20% through automation and optimized architectures.
- Leveraged Terraform to automate AWS infrastructure, increasing deployment speed by 40%, minimizing errors, and enhancing team agility.
- Utilized AWS CloudWatch and CloudTrail for security and compliance, implemented a cost optimization strategy, reducing bills by 23%, and maximizing ROI.
- Automated scripts with Ansible, reducing coding time by 17%, enhancing efficiency for rapid development and deployment processes.
- Crafted secure, resource-efficient Docker images with a 30% size reduction, orchestrating 80% faster deployments in Kubernetes clusters, maintaining 99.9% uptime.
- Built high-performance data pipelines with Confluent Platform, achieving a 20% increase in data throughput and a 30% reduction in latency. Devised "Julieops" for simplified Kafka operations, enhancing team agility and compliance.
- Designed data pipelines and dashboards for real-time visibility, reducing troubleshooting times by 30%.
- Designed a centralized logging cluster (ELK Stack) for 800+ containerized apps on OpenShift, improving debugging efficiency by 60% and proactively identifying issues.
- Proactively alerted on critical pod states and resource overconsumption with Bash scripts, reducing downtime by 25% and resource costs by 10%.
- Crafted custom scripts to safeguard OpenShift and Kubernetes cluster health with proactive alerts for node outages and resource overconsumption, achieving 99.5% uptime and 35% resource savings.
- Implemented DevSecOps practices, reducing vulnerabilities by 30% and securing code deployments by 40% in the CI/CD pipeline.
- Leveraged the Atlassian stack for optimized workflows, achieving a 15% improvement in developer productivity and a 10% reduction in project turnaround time.

PagerDuty

- Version Control: Git, Github, Bitbucket
- Application Management: Weblogic, Websphere, Tomcat

- Utilized AI-powered predictive analytics to forecast future resource needs, resulting in 25% better capacity planning and optimized resource allocation, minimizing cost overruns and downtime.
- Implemented AI tools for automated documentation generation and updates, streamlining the process by 30%, ensuring system changes are consistently documented and readily accessible, enhancing team productivity and knowledge sharing.

DevOps/Hadoop Engineer NTT DATA, Hyderabad | Client:T-Mobile, USA
11/2015 - 09/2017

- Implemented a Git branching strategy to mitigate code conflicts during releases, improving code integration efficiency by 25%.
- Increased the platform's capacity by 300% by architecting and deploying a highly available AWS infrastructure with 50 EC2 instances, 10 ELBs, and S3 storage, enabling smooth operation for 500 concurrent users.
- Managed S3 storage for backup, implementing efficient backup strategies and reducing data recovery time by 33%.
- Managed and optimized YARN, HDFS, MapReduce, and Hive clusters on 50+ nodes, ensuring 99.5% uptime and resource utilization of 85%.
- Improved Hadoop cluster performance by 25% through configuration optimization and job scheduling techniques.
- Architected and maintained Kafka clusters with 650+ topics and 500+ producers/consumers, handling 1 million events per second with an average latency of 5 milliseconds.
- Integrated Kafka with 5 internal databases and a cloud-based analytics platform for real-time decision-making.
- Created custom dashboards and alerts that reduced Kafka downtime by 30% through proactive issue identification.
- Automated routine administrative tasks using scripting languages, saving approximately 15 hours per week.
- Maintained comprehensive documentation for Hadoop and Kafka configurations and procedures.
- Also, conducted training sessions, improving team proficiency and knowledge sharing.
- Collaborated with vendors for issue resolution and performance optimization, resulting in a 20% faster turnaround time for critical issues.
- Proactively monitored and resolved issues in 12 production applications, ensuring 99.9% uptime and 50% reduction in support tickets.
- Engineered efficient data ingestion processes, optimizing the transfer of large datasets into Hadoop clusters, resulting in a 20% reduction in data loading times.

Systems Engineer Tivoli - Copenhagen
04/2011 - 08/2015

- Installed and configured Weblogic Server, ensuring optimal performance and reliability.
- Managed Weblogic domains, contributing to a 20% reduction in deployment time.
- Collaborated with database administrators to optimize server performance, resulting in a 15% improvement in database responsiveness.
- Monitored the health of Weblogic servers, reducing system downtime by

25% through proactive issue identification.

- Developed custom scripts for automation, streamlining routine tasks and reducing manual effort by 20%.
- Analyzed log files for performance issues, implementing corrective measures and reducing performance-related incidents by 30%.
- Handled SSL certificate installation/upgradation, ensuring secure communication and compliance with security standards.
- Managed incident requests, change requests, and problem statements, achieving a 15% reduction in issue resolution time.
- Conducted periodic backups and troubleshooting activities, ensuring data integrity and minimizing system disruptions.

EDUCATION

Completed coursework towards Master of Science (M.Sc) - Electronics Engineering
Osmania University - Hyderabad

CERTIFICATIONS

AZ 400: Designing and Implementing Microsoft DevOps Solutions

SAA-C03: AWS Certified Solutions Architect – Associate

CKA: Certified Kubernetes Administrator

CCAH: Cloudera Certified Administrator for Apache Hadoop

AWS Knowledge: Amazon EKS

ONSITE VISITS

Gained financial industry insights during a 2018 visit to Point72 New York with a valid US B1/132 VISA. Previously held a work visa for Denmark, showcasing international work experience and a commitment to diverse professional environments.